# STAT406- Methods of Statistical Learning Lecture 1

Matias Salibian-Barrera

UBC - Sep / Dec 2017

1

# About me

- Matias Salibian-Barrera
- matias@stat.ubc.ca
- http://www.stat.ubc.ca/~matias
- http://github.com/msalibian
- @msalibian
- Professor, Department of Statistics
- Undergrad in Math, PhD in Stats

2

# Prerequisites

- STAT306 `or` ECON326 `or` linear regression

- You are comfortable **working independently**

- You are **motivated** and enjoy being challenged

- You **want to be here** and are **interested in learning** the material

# Philosophy of the class

- We're here to **help you learn** (vs. teaching you)

- We'll encourage **engagement**, **curiosity** and **generosity**

- We'll have **zero tolerance for plagiarism**

- We favour **steady work** through the Term (vs. sleeping until finals)

4

# Lectures



- **Bring your laptop**
- **Prepare for class**
- Ask, doubt, question, discuss

5

# Lectures / Labs / Office hours

- Two weekly lectures, one weekly lab
- Ongoing evaluation - you are expected to attend **all** course meetings
- **Pre**-lecture readings and activities
- **Office hour**: Wed 1:30 - 2:30, ESB 3174
- It is a **4th year course** – expectations are high

# Grades

- Assignments: 10%,

- Quizzes (Webwork): 20%

- Lab activities: 10%

- In-Class midterm: 20%

- Final exam: 40%.

- There will be **no make-up** activities, quizzes, labs, assignments or exams. Anything you **miss** (with official documentation) will be assigned to your **final exam weight**.

# Textbook?

- No textbook

- Lecture slides + verbose scripts

- Several reference books – all available on-line @ UBC Library or publisher

- **Most used**: [**JWHT13**] - *An Introduction to Statistical Learning*, James, Witten, Hastie, Tibshirani, R., 2013, Springer, New York.

8

# Computer



- ## We will use R
  - Open source and free
  - Very flexible, relatively powerful
  - "Standard" in Statistics community, some industry
- ## Webwork quizzes rely on R

# Computer

- **YOU ARE RESPONSIBLE** for learning
- We **CAN HELP** with R
- We **WON'T teach all of** R
- There are tons of on-line resoures
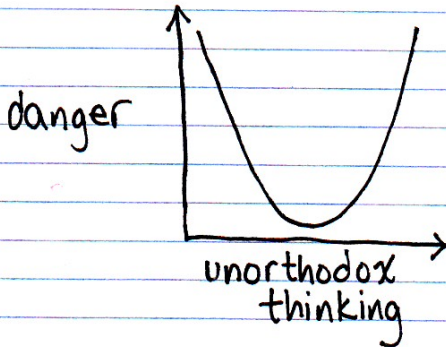
10

# Other resources

- Scripts to reproduce class examples on
  `github.com/msalibian/STAT406-2017`

- `PIAZZA`: **student led & driven** forum
  `piazza.com/ubc.ca/winterterm12017/stat406`

# Lectures?



thisisindexed.com

12

# Lectures?



thisisindexed.com

13
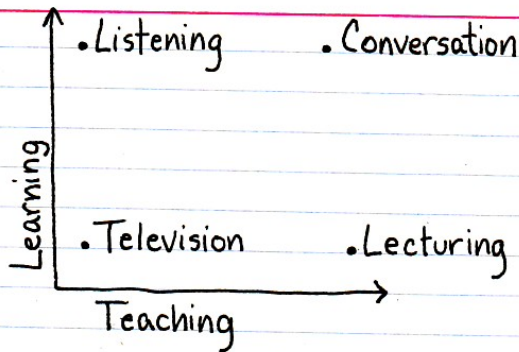
# Discussion

## Statistical learning

14

# Discussion

## Models versus "predictive algorithms"

# Review...

- $Y$ is the response variable
- **X** is a vector of auxiliary variables

$$Y = f(\mathbf{X}) + \varepsilon$$

- $f : \mathbb{R}^p \to \mathbb{R}$, unknown
- If $E[\varepsilon] = 0$

$$E[Y|\mathbf{X}] = f(\mathbf{X})$$

16

# Review...

- In a linear model, $f$ is a linear function

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \cdots + \beta_p X_p + \varepsilon$$

- If $E[\varepsilon] = 0$

$$E[Y|X_1, X_2, \ldots, X_p] = \\ \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \cdots + \beta_p X_p$$

# Review...

- Why would we want to estimate the coefficients of the linear model?

- What's the connection with prediction?

# Review...

- Why would we want to estimate the coefficients of the linear model?

- What's the connection with prediction?

- "Best predictor"

$$\arg \min_{\mathbf{h}} E\left[(Y - \mathbf{h}(\mathbf{X}))^2\right] = E[Y|\mathbf{X}]$$

# Review...

- Best predictor is the regression function

- We need to estimate $E[Y|\mathbf{X}]$

- We propose a model (e.g. linear) for $E[Y|\mathbf{X}]$ and estimate it

- E.g. in a linear model, to estimate $f(\mathbf{X})$ we need to estimate $\beta_0, \beta_1, \ldots, \beta_p$

19

# Review...

- Data $(Y_1, \mathbf{X}_1), (Y_2, \mathbf{X}_2), \ldots, (Y_n, \mathbf{X}_n)$,

- Least squares estimator

$$\hat{\boldsymbol{\beta}} = \arg\min_{\beta_0, \boldsymbol{\beta}} \sum_{i=1}^{n} \left(Y_i - \beta_0 - \boldsymbol{\beta}' \mathbf{X}_i\right)^2$$

20

# Review...

- There is a closed form for $\hat{\beta}$

$$\hat{\beta} = (X'X)^{-1} X'\mathbf{Y}$$

where $\mathbf{Y} = (Y_1, Y_2, \ldots, Y_n)'$ and

$$X = \begin{pmatrix} 1 & \cdots & \mathbf{X}_1 & \cdots \\ 1 & \cdots & \mathbf{X}_2 & \cdots \\ 1 & \cdots & \mathbf{X}_3 & \cdots \\ \vdots & \cdots & \cdots & \cdots \\ 1 & \cdots & \mathbf{X}_n & \cdots \end{pmatrix}$$

21

# Review...

- As long as $E[\varepsilon] = E[\varepsilon|\mathbf{X}] = 0$ we have

$$E\left[\hat{\beta}\right] = \beta = \begin{pmatrix} \beta_0 \\ \beta_1 \\ \vdots \\ \beta_p \end{pmatrix}$$

- The LS estimator is consistent and unbiased

- Do we need any other assumption?

# Review...

- Consider the air pollution data
- $n = 60$ observations
- $p = 16$, response variable: MORT
- A linear model:

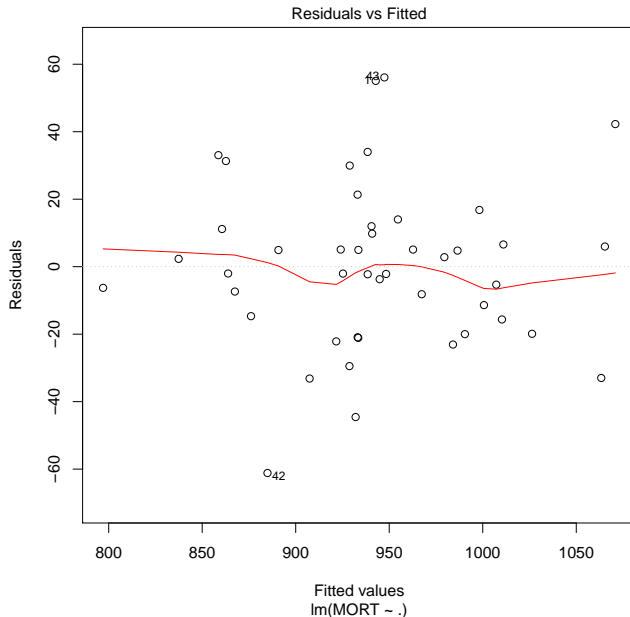$$\text{MORT} = \beta_0 + \beta_1 \, \text{PREC} + \beta_2 \, \text{JANT} + \ldots + \epsilon$$

or equivalently

$$E\left(\text{MORT} \,\middle|\, \text{PREC}, \text{JANT}, \ldots\right)$$
$$= \beta_0 + \beta_1 \, \text{PREC} + \beta_2 \, \text{JANT} + \ldots$$

# Review...

- Randomly split into a training ($n=45$) and a test set ($n=15$)
- Use training set to fit a model

- Read data into object `x.tr`
- Fit the "full" model
- "Look" at the fit

# Diagnostics



Residuals vs Fitted

25

# Diagnostics



Normal Q–Q

# Diagnostics

```
> full <- lm(MORT ~ ., data=x.tr)
> summary(full)

> sum( resid(full)^2 )
[1] 25898.8
```
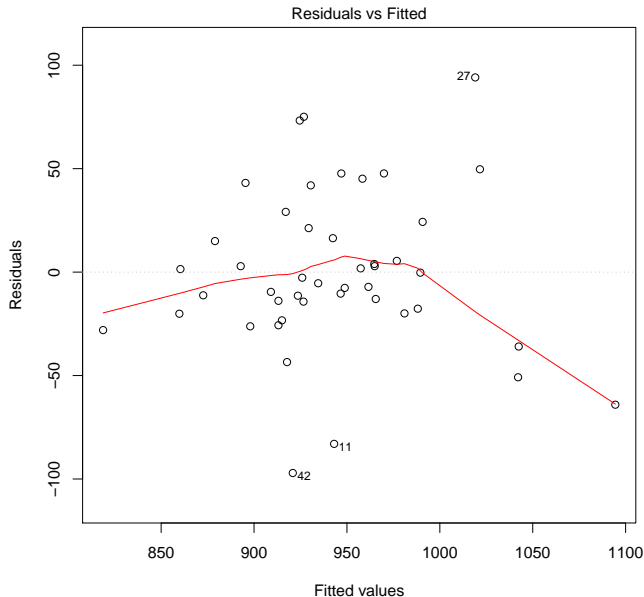
# Fit a reduced model

```
> reduced <- lm(MORT ~ POOR + HC +
NOX + HOUS + NONW, data=x.tr)

> sum( resid(reduced)^2 )
[1] 66135.29
```

# Diagnostics for reduced model



Residuals vs Fitted

lm(MORT ~ POOR + HC + NOX + HOUS + NONW)

29

# Diagnostics for reduced model



Normal Q–Q

lm(MORT ~ POOR + HC + NOX + HOUS + NONW)

30

# Discussion

## Goodness of fit versus prediction power

# Predictions

$$Y \longleftrightarrow \hat{f}(\mathbf{X})$$

$$\left(Y - \hat{f}(\mathbf{X})\right)^2 \quad ?$$

$$E\left[\left(Y - \hat{f}(\mathbf{X})\right)^2\right] \quad ?$$

What are we "averaging" over?
What is random?

32

# Predictions

$$E_{(Y^*, \mathbf{x}^*)} \left[ \left( Y^* - \hat{f}(\mathbf{X}^*) \right)^2 \right]$$

where $(Y^*, \mathbf{X}^*)$ are new, future observations, not used when computing $\hat{f}$.

# Predictions

If we assume that $Y = f(X) + \epsilon$, then

$$E_{(Y^*, \mathbf{X}^*)}\left[\left(Y^* - \hat{f}(\mathbf{X}^*)\right)^2\right] =$$

$$E_{(Y^*, \mathbf{X}^*)}\left[\left(f(\mathbf{X}^*) - \hat{f}(\mathbf{X}^*)\right)^2\right] + V(\epsilon)$$

- what assumptions are needed for this to be true?
- is it still true if I look at predictions for a single & fixed $\mathbf{X}_0$?

# Predictions

- What we want

$$E_{(Y^*, \mathbf{x}^*)}\left[\left(Y^* - \hat{f}(\mathbf{X}^*)\right)^2\right]$$

  is very difficult to estimate

- Something similar

$$E_{\{(Y^*, \mathbf{x}^*), \text{data}\}}\left[\left(Y^* - \hat{f}(\mathbf{X}^*)\right)^2\right]$$

  is easier to estimate

# Predictions

- Read the test set

- Use both models to predict `MORT`

- Compare both sets of predictions

# Predictions

```
> x.te <- read.table('pollution-test.dat',...
>
> x.te$pr.full <- predict(full, newdata=x.te)
> x.te$pr.reduced <- predict(reduced,
        newdata=x.te)
>
> with(x.te, mean( (MORT - pr.full)^2 ))
[1] 4677.45
>
> with(x.te, mean( (MORT - pr.reduced)^2 ))
[1] 1401.571
```

## Discuss

# Discussion points

- Goodness of fit vs. prediction power

- How do we estimate prediction MSE?

$$E_{(Y^*, \mathbf{x}^*)} \left[ \left( Y^* - \hat{f}(\mathbf{X}^*) \right)^2 \right]$$

- Can it be done without a test set?

# Next week...

- Quiz 0 (Review) is out, due next class.
- Check `connect.ubc.ca` often
- Visit `github.com/msalibian/STAT406-2017`
- Use `PIAZZA`
- Read the suggested sections of [JWHT13]
- Attend the lab